

PracticalMachineLearningFinalProject

Wanhao Chi

October 13, 2016

Data Process

1. Read the datasets and do data cleaning

```
### read data
training <- read.csv(url("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"),na.strings = c("NA","#div/0!",""))
testing <- read.csv(url("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"),na.strings = c("NA","#div/0!",""))
### take a first look at the data
# dimersions of the dataset
dim(training);dim(testing)
```

```
## [1] 19622 160
```

```
## [1] 20 160
```

```

#list types for each attribute
#sapply(training,class)
# how many types in total
#table(sapply(training,class));table(sapply(testing,class))
# take a look at the first 2 rows of the data
#head(training, n = 2);head(testing, n =2)
# summarize attribute distributions
#summary(training); summary(testing)
### data cleaning
# remove the first 7 columns that have nothing with the predictions
training <- training[,-c(1:7)]
testing <- testing[,-c(1:7)]
# remove the columns that contain missing values
training <- training[,colSums(is.na(training)) == 0]
testing <- testing[,colSums(is.na(testing)) == 0]

# take a look at the data again
# dimersions of the dataset
dim(training);dim(testing)

```

```
## [1] 19622    53
```

```
## [1] 20 53
```

2. Create a validation dataset from the Training dataset

```

set.seed(123)
library(caret)

```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```

validation_index <- createDataPartition(training$classe, p =0.8, list = FALSE)
validation <- training[-validation_index,]
training <- training[validation_index,]

```

Model selection

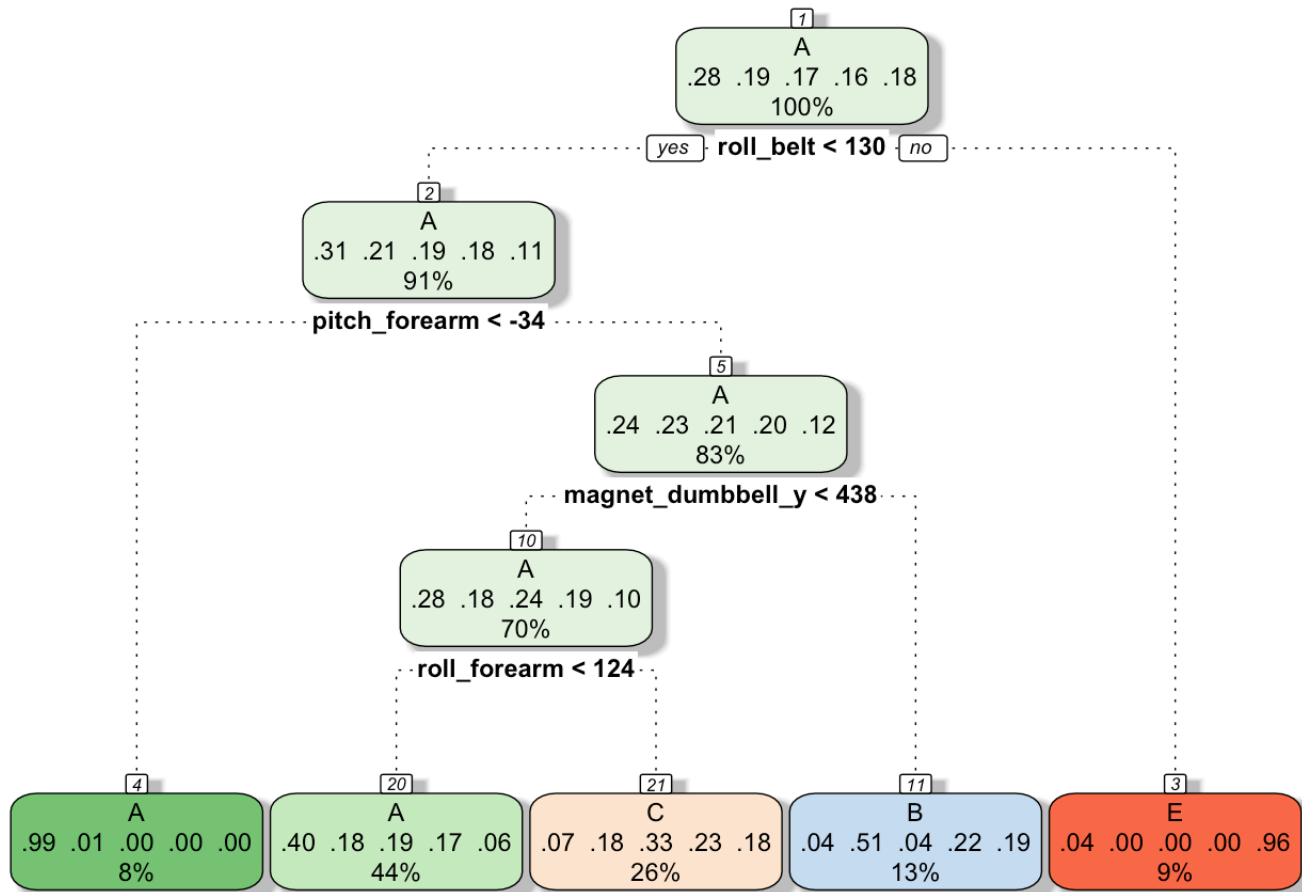
Since the outcome is categorical, we will use classification trees and random forests to predict it.

1. Classification trees

```
set.seed(125)
ctrl <- trainControl(method = "cv", number = 10) # 10 fold cross-validation is used in general
modFit_rpart <- train(classe ~ ., data= training, method = "rpart", trControl = ctrl, na.action = na.omit)
print(modFit_rpart, digits = 3)
```

```
## CART
##
## 15699 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 14129, 14127, 14129, 14130, 14129, 14130, ...
## Resampling results across tuning parameters:
##
##    cp      Accuracy  Kappa
##    0.0368  0.502     0.3489
##    0.0601  0.414     0.2063
##    0.1170  0.332     0.0734
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0368.
```

```
library("rattle")
fancyRpartPlot(modFit_rpart$finalModel)
```



Rattle 2016-Oct-13 20:58:31 wanhaochi

```

# predict outcomes using validation dataset
predict_rpart <- predict(modFit_rpart, validation)
confusion_rpart <- confusionMatrix(validation$classe, predict_rpart)
confusion_rpart

```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
##           A 987   14   93    0   22
##           B 312  271  176    0    0
##           C 317   29  338    0    0
##           D 296  125  222    0    0
##           E 108   80  200    0  333
##
## Overall Statistics
##
##           Accuracy : 0.4917
##           95% CI : (0.476, 0.5075)
##           No Information Rate : 0.5149
##           P-Value [Acc > NIR] : 0.9983
##
##           Kappa : 0.3361
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.4886  0.52216  0.32847      NA  0.93803
## Specificity      0.9322  0.85664  0.88044  0.8361  0.89126
## Pos Pred Value   0.8844  0.35705  0.49415      NA  0.46186
## Neg Pred Value   0.6320  0.92162  0.78666      NA  0.99313
## Prevalence       0.5149  0.13230  0.26230  0.0000  0.09049
## Detection Rate   0.2516  0.06908  0.08616  0.0000  0.08488
## Detection Prevalence 0.2845  0.19347  0.17436  0.1639  0.18379
## Balanced Accuracy 0.7104  0.68940  0.60446      NA  0.91464
```

```
accuracy_rpart <- confusion_rpart$overall[1]
accuracy_rpart
```

```
## Accuracy
## 0.4917155
```

Therefore, the accuracy using classification tree is 0.492. The out-of-sample error rate is 0.508.

2. random forest

```
set.seed(124)
#library("randomForest")
ctrl <- trainControl(method = "cv", number = 5) # 5 fold cross-validation is used
modFit_rf <- train(classe ~ ., data= training, method = "rf",trControl = ctrl,na.acti
on = na.omit)
#modFit_rf <- randomForest(classe ~ ., data = training)
print(modFit_rf, digits = 3)
```

```
## Random Forest
##
## 15699 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 12558, 12560, 12562, 12558, 12558
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa
##    2    0.991    0.989
##   27    0.992    0.990
##   52    0.985    0.981
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was mtry = 27.
```

```
# predict outcomes using validation dataset
predict_rf <- predict(modFit_rf, validation)
confusion_rf <- confusionMatrix(validation$classe, predict_rf)
confusion_rf
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A     B     C     D     E
##           A 1115     1     0     0     0
##           B   0   758     1     0     0
##           C   0     1   683     0     0
##           D   0     0    7   636     0
##           E   0     0     0     0   721
##
## Overall Statistics
##
##           Accuracy : 0.9975
##           95% CI : (0.9953, 0.9988)
##           No Information Rate : 0.2842
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9968
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           1.0000    0.9974    0.9884    1.0000    1.0000
## Specificity           0.9996    0.9997    0.9997    0.9979    1.0000
## Pos Pred Value        0.9991    0.9987    0.9985    0.9891    1.0000
## Neg Pred Value        1.0000    0.9994    0.9975    1.0000    1.0000
## Prevalence            0.2842    0.1937    0.1761    0.1621    0.1838
## Detection Rate        0.2842    0.1932    0.1741    0.1621    0.1838
## Detection Prevalence  0.2845    0.1935    0.1744    0.1639    0.1838
## Balanced Accuracy      0.9998    0.9985    0.9941    0.9989    1.0000
```

```
accuracy_rf <- confusion_rf$overall[1]
accuracy_rf
```

```
## Accuracy
## 0.9974509
```

Therefore, the accuracy using random forest is 0.997, which is much better than what we got from classification tree method. The out-of-sample error for random forest method is 0.003. However, the random forest method is computationally inefficient.

Prediction on testing dataset

```
prediction_testing <- predict(modFit_rf,testing)
prediction_testing
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```